

COMPSCI 250: Introduction to Computation

Lecture #37: Two-Way Automata and Turing Machines
David Mix Barrington and Ghazaleh Parvini
12 May 2023

2WDFA's and Turing Machines

- Enhancing a DFA's Abilities
- Definition and Semantics of 2WDFA's
- Why 2WDFA's Have Regular Languages (Sketch)
- Turing Machines
- The Formal Turing Machine Model
- A Turing Machine Example
- The Church-Turing Thesis

Enhancing a DFA's Abilities

- DFA's, and the other models we have now shown to be equivalent to them, model a particular kind of computation. A DFA:
- (1) can read its input only once, from left to right,
- (2) can only read, not write to, the memory holding the input, and
- (3) has only a bounded amount of memory apart from that input.

Enhancing a DFA's Abilities

- In our last week of lectures we will look at another model of computation called a **Turing machine**, which we can think of as an enhanced DFA. Turing machines:
- (1) can move both ways on the **tape** that contains their input,
- (2) can **write** new characters into the space that originally holds the input, and
- (3) can utilize **additional memory**, as much as they need, as well as the original space.

Enhancing a DFA's Abilities

- We'll begin today by looking at the effect of adding new ability (1) alone to a DFA, producing a new kind of machine called a **two-way DFA**.
- In COMPSCI 501 you'll also look at machines that have new abilities (1) and (2) but not (3) -- these are called **linear bounded automata**.

Two-Way Finite Automata

- Like a DFA, a 2WDFA has a state set Q , start state i , final state set F , input alphabet Σ , and transition function δ .
- The only difference is that δ goes from $Q \times \Sigma$ to $Q \times \{L, R\}$. Based on the current state and the letter it sees, the 2WDFA enters a new state and moves *either left or right* on its tape.
- It continues taking steps until or unless it moves off one end of the tape.

Semantics of 2WDFA's

- We need to define the **semantics** of the 2WDFA M -- the meaning of each computation in terms of defining a language $L(M)$.
- We start with the **read head** on the first letter of the input, and start the computation.
If the machine moves off the left end of the tape, we say that it **hangs** and the input is not in $L(M)$.

Semantics of 2W DFA's

- If it moves off the right end of the tape, we say that it **accepts** if it goes into a final state and that it **rejects** if it goes into a nonfinal state.
- There is a fourth possibility, that it **loops** or never terminates.
- The input is in $L(M)$ if and only if M accepts.


A 2WDFA Example

- Let's look at the behavior of this 2WDFA on some strings:

- On a, it moves right off the input in state f and accepts. iRf

- On b, it moves off the left end and hangs.


- On aaa, it moves right to state f, right again to state p, left to state f, right to p,..., and thus loops forever. iLp hang



iRfRpLfRpLf...

Clicker Question #1


- What does this 2WDFA do on input string babab?
- (a) accepts by leaving to the right in a final state
- (b) enters an infinite loop
- (c) hangs by leaving to the left
- (d) rejects by leaving to the right in a nonfinal state



Not the Answer

Clicker Answer #1

- What does this 2WDFA do on input string babab?
- (a) accepts by leaving to the right in a final state
- *(b) enters an infinite loop*
- (c) hangs by leaving to the left
- (d) rejects by leaving to the right in a nonfinal state



iRpRfRiLfRiLfR...
b₁a₂b₃a₄b₃a₄b₃...

2WDFA's and Regular Languages

- Could a 2WDFA have a non-regular language like $\{a^n b^n : n \geq 0\}$?

For DFA's, we argued that after the a's have been read, the machine “must know” how many a's it saw (formally, each different number of a's was in a different equivalence class).

- But now, the machine could make multiple visits to the a's. Can it use this capability in any way to get more information about the a's?



2WDFA's and Regular Languages

- In Section 15.1 of the text, we prove that the language of any 2WDFA is regular. Here is a sketch of the argument.
- Given a 2WDFA M and a string w , we define several functions of w based on M 's behavior.
- If M exits w to the right in state q when started in state i on the left, we say that $f_0(w) = q$.
- If it hangs or loops in that situation, we say that $f_0(w) = d$.

2WDFA's and Regular Languages

- Similarly, we define a function f_p for each state p . Consider starting M on the *right* of w in state p .
- If it loops or hangs, we define $f_p(w) = d$.
- If it exits to the right in state q , we define $f_p(w) = q$.

Converting a 2WDFA Example



Language of DFA = $a(ba + aba)^*$

2W DFA's = Regular

- Here's the crux of the argument. Suppose that for two strings v and w , the values of each of these functions are the same. That is, $f_0(v) = f_0(w)$ and for each state p , $f_p(v) = f_p(w)$.
- Then, we will argue, v and w are $L(M)$ -equivalent in the sense of the Myhill-Nerode Theorem.
- Since there are only finitely many possible sequences of values for these functions, there are only finitely many equivalence classes, and the theorem tells us that $L(M)$ is a regular language.

2WDFA's = Regular

- We need to show that for any string z , the strings vz and wz are either both in $L(M)$ or both not in $L(M)$.
- Let z be an arbitrary string, assume that the functions agree on v and w , and look at what happens when M starts computing on v and on w .
- If M hangs or loops on vz without leaving v , it must do the same on wz because $f_0(v) = f_0(w) = d$.

2WDFA's = Regular

- If it exits v to the right, then it also exits w to the right, and in the same state. From that point, the two computations in z proceed identically, until or unless they leave z .
- If they leave to the right, both computations accept or both reject. If they go back into v and w , they do so in the same state p . Then either both die, or both move back into z in the same state $f_p(v) = f_p(w)$, and so forth until eventually both accept, both reject, or both die. So $vz \in L(M) \leftrightarrow wz \in L(M)$.

Turing Machines

- In the 1930's, various researchers designed **systems of computation** in an attempt to create a simple mathematically precise model that could express any possible computation. The model that has become most widely used is the **Turing machine**, proposed by the English mathematician Alan Turing in 1936. (Another one of these models, the **lambda calculus** of Alonzo Church, developed into the Lisp family of programming languages.)

Turing Machines

- Turing and Church each convinced themselves that any clear, precise computational instructions could be translated (we might say “compiled”) into each of their systems.
- When each heard about the other’s system, they proved that any computation in one could be translated to the other.

Turing Machines


- Thus the two systems defined the same set of **computable functions** from strings to strings.
- Just as a language is either regular or not, a function is either computable or not.
- (Actually finite-state machines would not be formalized for another twenty years or so.)

The Turing Machine Model

- A **Turing machine** is formally defined by giving a state set Q , an input alphabet Σ , a start state i , and a final state set F , as we've seen already.
- But it also has a **tape alphabet** Γ with $\Sigma \subseteq \Gamma$, and a **blank symbol** \square that is an element of Γ and is the initial contents of every tape cell right of the input.

The Turing Machine Model

- The machine has a **tape** that is infinite to the right and finite to the left. Each cell of the tape holds a letter in Γ at any given time.
- There is a **head** that points to one cell of the tape at any given time.



The Turing Machine Model


- The **transition function** δ is from $Q \times \Gamma$ to $Q \times \Gamma \times \{L, R\}$.

A **step** of the computation consists of the machine looking at the letter at its head, applying δ to its current state and that letter to get a triple (q, a, L) or (q, a, R) , then *changing its state* to q , *writing* an a in the current cell, and *moving* left or right.

- Actually δ is not defined for states in F -- the machine **halts** in those states.

Turing Machine Configurations


- At any given time, we can describe everything we would ever want to know about the Turing machine's computation by a string called a **configuration**.
- What we need to record is the current state, the contents of the tape, and the position of the head.



Configuration: □aaqb□bab□

Turing Machine Configurations


- We record the tape contents as a string of letters from Γ , starting at the left end of the tape and ending with the last non-blank letter.
- We record the state and head position by inserting a letter for the state into this string, just to the left of the head position.



Configuration: □aaqb□bab□

Clicker Question #2

- Suppose $\delta(q, b) = (r, \square, L)$. What will be the new configuration of the Turing machine below?
- (a) $\square baa\square bbb\square$ (b) $\square bar\square\square bbb\square$
- (c) $\square bra\square\square bbb\square$ (d) $\square babr\square bbb\square$




Configuration: $\square baqb\square bbb\square$

Not the Answer

Clicker Answer #2

- Suppose $\delta(q, b) = (r, \square, L)$. What will be the new configuration of the Turing machine below?
- (a) $\square baa\square bbb\square$ (b) $\square bar\square\square bbb\square$
- (c) $\square bra\square\square bbb\square$ (d) $\square babr\square bbb\square$



Configuration: $\square baqb\square bbb\square$

Turing Machine Configurations

- A Turing machine starts with only finitely many non-blank symbols on its tape.
- So in writing a configuration, we only need to go to the last non-blank symbol (unless we need to go further to indicate the head position).
- We can think of the computation then as a series of configurations, starting with $i \sqcap w_1 w_2 \dots w_n$ and continuing until or unless the machine halts or hangs.

A Turing Machine Example

- On the next slide is a machine that solves a problem that a DFA cannot. When started in configuration $i \sqsubset w_1 w_2 \dots w_n$, it will halt if and only if w is in the language $\{a^n b^n : n \geq 0\}$ -- otherwise it will hang.
- With input $aabb$ we get $i \sqsubset aabb$, $\sqsubset paabb$, $\sqsubset \sqsubset qabb$, $\sqsubset \sqsubset aqbb$, $\sqsubset \sqsubset abqb$, $\sqsubset \sqsubset abbq \sqsubset$, $\sqsubset \sqsubset abrb$, $\sqsubset \sqsubset asb$, $\sqsubset \sqsubset sab$, $\sqsubset s \sqsubset ab$, $\sqsubset \sqsubset pab$, $\sqsubset \sqsubset \sqsubset qb$, $\sqsubset \sqsubset \sqsubset bq \sqsubset$, $\sqsubset \sqsubset \sqsubset rb$, $\sqsubset \sqsubset s \sqsubset$, $\sqsubset \sqsubset \sqsubset p \sqsubset$, $\sqsubset \sqsubset \sqsubset h \sqsubset$. The string $aabb$ is accepted.

A Turing Machine Example

In i: Move R and go to p.

In p: On \square , go to h.

On b, move L and go to z.

On a, print \square , move R, and go to q.

In q: On a or b, move R and stay in q.

On \square , move L and go to r.

In r: On a or \square , move L and go to z.

On b, print \square , move L, and go to s.

In s: On a or b, move L and stay in s.

On \square , move R and go to p.

In h: Halt (final state).

In z: Move left and stay in z.

Clicker Question #3

- What does this TM do on inputs aa (starting i□aa□) and bb (starting i□ab□)?
- (a) accepts both (b) accepts aa, hangs on ab
- (c) hangs on both (d) hangs on aa, accepts ab

In i: Move R and go to p.

In p: On □, go to h.

On b, move L and go to z.

On a, print □, move R, and go to q.

In q: On a or b, move R and stay in q.

On □, move L and go to r.

In r: On a or □, move L and go to z.

On b, print □, move L, and go to s.

In s: On a or b, move L and stay in s.

On □, move R and go to p.

In h: Halt (final state).

In z: Move left and stay in z.

Not the Answer

Clicker Answer #3

- What does this TM do on inputs aa (starting i□aa□) and bb (starting i□ab□)?
- (a) accepts both (b) accepts aa, hangs on ab
- (c) hangs on both *(d) hangs on aa, accepts ab*

In i: Move R and go to p.

In p: On □, go to h.

On b, move L and go to z.

On a, print □, move R, and go to q.

In q: On a or b, move R and stay in q.

On □, move L and go to r.

In r: On a or □, move L and go to z.

On b, print □, move L, and go to s.

In s: On a or b, move L and stay in s.

On □, move R and go to p.

In h: Halt (final state).

In z: Move left and stay in z.

i□aa□

□paa□

□□qa□

□□aq

□□ra□

□z□a□

z□□a□

hang

i□ab□

□pab□

□□qb□

□□bq□

□□rb□

□s□

□□p□

□□h□□

accept

The Church-Turing Thesis

- The **Church-Turing Thesis** says that any “reasonable” general-purpose model of computation will be able to compute exactly the same functions from strings to strings as Turing machines or the lambda calculus.
- (More precisely, they compute the same set of **partial functions**, because a general computation always has the possibility of not returning an output.)

The Church-Turing Thesis

- We can't mathematically prove this thesis, only amass evidence for it. In fact it actually serves as an implicit definition of "reasonable".
- Serious people have argued against the thesis -- for example physicist Roger Penrose argues that quantum effects in the brain compute in ways that a Turing machine could not. (Dave doesn't buy it.)
- For more on this see Turing's article *On Minds and Machines* or almost anything by Douglas Hofstadter.

The Church-Turing Thesis

- You probably believe that we could simulate a Turing machine in Java, given unlimited memory. Could a Turing machine simulate any Java program?
- We know that Java can be compiled into machine language, so we would have to believe that any machine language program could be simulated by a TM.